

Chapter 6

Heapsort

Chi-Yeh Chen

陳奇業

成功大學資訊工程學系



藏行顯光
成就共好

Achieve Securely
Prosper Mutually



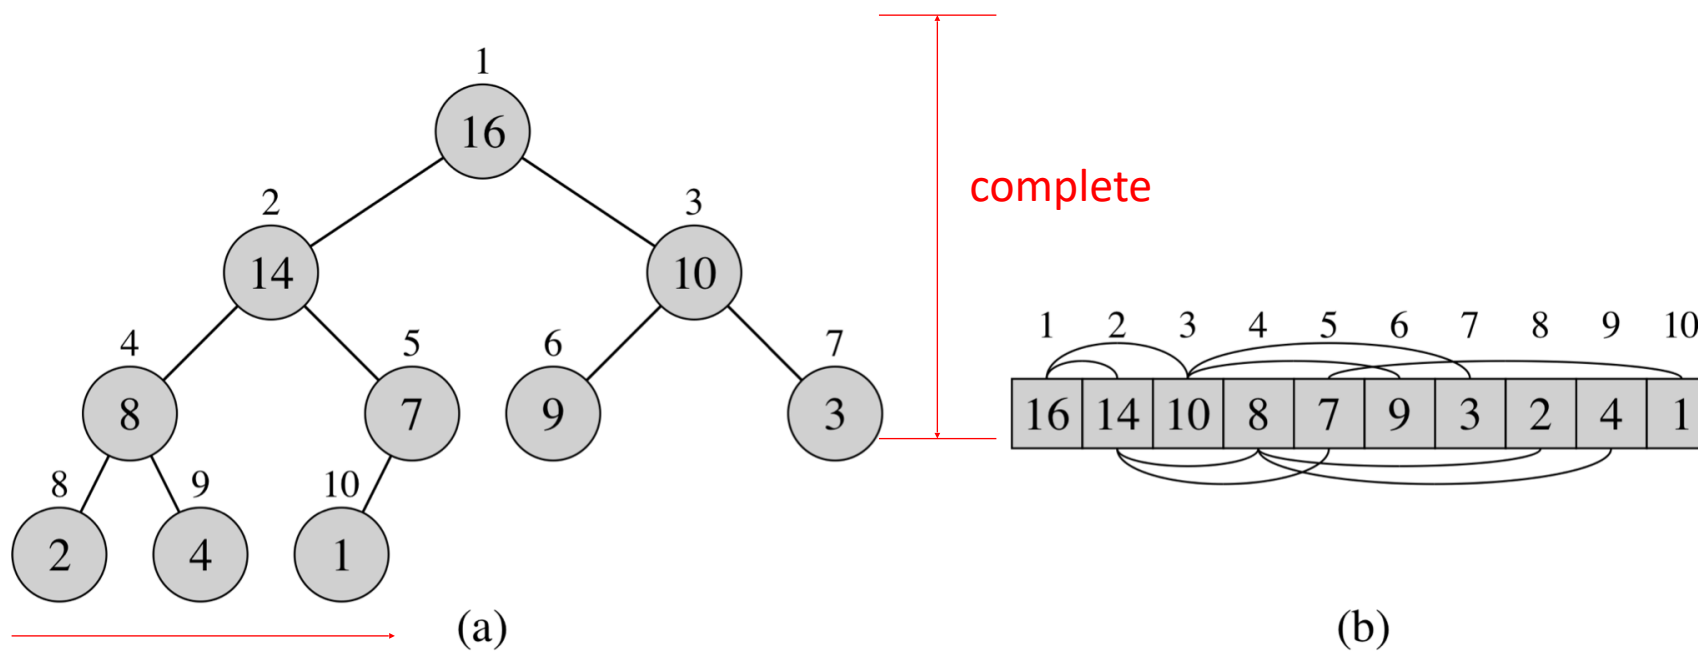
國立成功大學 九十週年
90th Anniversary of NCKU

Overview

- Like merge sort, but unlike insertion sort, heapsort's running time is $O(n \lg n)$.
- Like insertion sort, but unlike merge sort, heapsort sorts in place.



Heap



from the left up to a point

Heap

- Heap A is a nearly complete binary tree.
 - **Height** of node = # of edges on a longest simple path from the node down to a leaf.
 - **Height** of heap = height of root = $\Theta(\lg n)$



Heap

- $A.length$, which gives the number of elements in the array
- $A.heap_size$, which represents how many elements in the heap are stored within array A .
- Although $A[1 \dots A.length]$ may contain numbers, only the elements in $A[1 \dots A.heap_size]$

Heap

- A heap can be stores as an array A
 - **Root** of trees is $A[1]$
 - **Parent** of $A[i] = A[\lfloor i/2 \rfloor]$
 - **Left child** of $A[i] = A[2i]$
 - **Right child** of $A[i] = A[2i + 1]$
 - Computing is fast with binary representation implementation

Heap

Parent(i)

1 return $\lfloor i/2 \rfloor$

LEFT(i)

1 return $2i$

RIGHT(i)

1 return $2i + 1$




Heap property

- Heap property
 - For max-heap (largest element at root), *max-heap property*: for all nodes i , excluding the root, $A[PARENT(i)] \geq A[i]$.
 - For min-heap (smallest element at root), *min-heap property*: for all nodes i , excluding the root, $A[PARENT(i)] \leq A[i]$.
- Maximum element of a max-heap is at the root.
- The heapsort algorithm we'll use max-heaps.

Example

- A max-heap



Heap property

- The basic operations on heaps run in time at most proportional to the height of the tree and thus take $O(\lg n)$ time.
 - MAX-HEAPIFY: $O(\lg n)$
 - BUILD-MAX-HEAP: run in linear time $O(n)$.
 - HEAPSORT: $O(n \lg n)$
 - MAX-HEAP-INSERT, HEAP-EXTRACT-MAX, HEAP-INCREASE-KEY and HEAP-MAXIMUM: $O(\lg n)$

- A heap with height h is an almost-complete binary tree (complete at all levels except possibly the lowest)
 - at most $2^{h+1} - 1$ elements (if it is complete)
 - at least $2^h - 1 + 1 = 2^h$ elements (if the lowest level has just 1 element and the other levels are complete).
- Given an n -element heap of height h , we know that $2^h \leq n \leq 2^{h+1} - 1 < 2^{h+1}$. Thus, $h \leq \lg n < h + 1$. Since h is an integer, $h = \lfloor \lg n \rfloor$.

Maintaining the heap property

- **MAX-HEAPIFY** is important for manipulating max-heaps. It is used to **maintain the max-heap property**.
 - Before MAX-HEAPIFY, $A[i]$ may be smaller than its children.
 - Assume left and right subtrees of i are max-heaps.
 - After MAX-HEAPIFY, subtree rooted at i is a max-heap



Pseudocode




MAXHEAPIFY(A, i, n)

```
1  $l \leftarrow LEFT(i)$ 
2  $r \leftarrow RIGHT(i)$ 
3 if  $l \leq n$  and  $A[l] > A[i]$  then
4    $largest \leftarrow l$ 
5 else
6    $largest \leftarrow i$ 
7 if  $r \leq n$  and  $A[r] > A[largest]$  then
8    $largest \leftarrow r$ 
9 if  $largest \neq i$  then
10   exchange  $A[i] \leftrightarrow A[largest]$ 
11   MAXHEAPIFY( $A, largest, n$ )
```




Example

- Run MAX-HEAPIFY on the following heap example.



Analysis of MAX-HEAPIFY



The children's subtrees each have size at most $2n/3$

Analysis of MAX-HEAPIFY

- $T(n) \leq T\left(\frac{2n}{3}\right) + \Theta(1)$

By case 2 of the master theorem, the solution to this recurrence is

$$T(n) = O(\lg n).$$

- Case 2: If $f(n) = \Theta(n^{\log_b a})$, then $T(n) = \Theta(n^{\log_b a} \lg n)$

Analysis of MAX-HEAPIFY

- **Time:** $O(\lg n)$
- **Correctness:** Heap is almost-complete binary tree, hence must process $O(\lg n)$ levels, with constant work at each level (comparing 3 items and maybe swapping 2).



- The following procedure, given an unordered array, will produce a max-heap.

BUILD-MAX-HEAP(A, n)

```
1 for  $i = \lfloor n/2 \rfloor$  downto 1 do  
2   MAX-HEAPIFY( $A, i, n$ )
```

$\lfloor n/2 \rfloor + 1, \lfloor n/2 \rfloor + 2, \dots, n$.

Assume $A[\lfloor n/2 \rfloor + 1]$ is not a leaf, then $\text{LEFT}(\lfloor n/2 \rfloor + 1) \leq n$





$\text{LEFT}(\lfloor n/2 \rfloor + 1) = 2(\lfloor n/2 \rfloor + 1) > 2(n/2 - 1 + 1) = n \rightarrow \leftarrow$



Example









- Building a max-heap from the following unsorted array results in the first heap example.
 - i starts off as 5.
 - MAX-HEAPIFY is applied to subtrees rooted at nodes (in order): 16, 2, 3, 1, 4.

	1	2	3	4	5	6	7	8	9	10
A	4	1	3	2	16	9	10	14	8	7

	1	2	3	4	5	6	7	8	9	10
A	4	1	3	2	16	9	10	14	8	7







Correctness

- A **loop invariant** is a property of a program loop that is true before (and after) each iteration.
- We must show **three things** about a loop invariant:
 - **Initialization**: It is true prior to the first iteration of the loop.
 - **Maintenance**: If it is true before an iteration of the loop, it remains true before the next iteration.
 - **Termination**: When the loop terminates, the invariant gives us a useful property that helps show that the algorithm is correct.

Correctness

- Loop invariant: At start of every iteration of for loop, each node $i + 1, i + 2, \dots, n$ is root of a max-heap.
 - Initialization: We know that each node $\lfloor n/2 \rfloor + 1, \lfloor n/2 \rfloor + 2, \dots, n$ is a leaf, which is the root of a trivial max-heap. Since $i = \lfloor n/2 \rfloor$ before the first iteration of the for loop, the invariant is initially true.
 - Maintenance: Children of node i are indexed higher than i , so by the loop invariant, they are both roots of max-heaps. Correctly assuming that $i + 1, i + 2, \dots, n$ are all roots of max-heaps, MAX-HEAPIFY makes node i a max-heap root. Decrementing i reestablishes the loop invariant at each iteration.
 - Termination: When $i = 0$, the loop terminates. By the loop invariant, each node, notably node 1, is the root of a max-heap.

- Be careful not to confuse the height of a node (longest distance from a leaf) with its depth (distance from the root).
- If the heap is not a complete binary tree (bottom level is not full), then the nodes at a given level (depth) don't all have the same height. For example, although all nodes at depth H have height 0, nodes at depth $H-1$ can have either height 0 or height 1.

Analysis

- There are at most $\lceil n/2^{h+1} \rceil$ nodes of height h in any n -element heap.
 - For a complete binary tree, it's easy to show that there are $\lceil n/2^{h+1} \rceil$ nodes of height h .
 - But the proof for an incomplete tree is tricky and is not derived from the proof for a complete tree.

Proof. By induction on h . Let H be the height of the heap.

Basis: Show that it's true for $h = 0$ (i.e., that # of leaves $\leq \lceil n/2^{h+1} \rceil = \lceil n/2 \rceil$). In fact, we'll show that the # of leaves = $\lceil n/2 \rceil$.

The tree leaves (nodes at height 0) are at depths H and $H - 1$. They consist of

- all nodes at depth H , and
- the nodes at depth $H - 1$ that are not parents of depth- H nodes.

Let x be the number of nodes at depth H —that is, the number of nodes in the bottom (possibly incomplete) level.

Note that $n - x$ is odd, because the $n - x$ nodes above the bottom level form a complete binary tree, and a complete binary tree has an odd number of nodes. Thus if n is odd, x is even, and if n is even, x is odd.

- If n is odd, then x is even, so all nodes have siblings—i.e., all internal nodes have 2 children. Thus (see Exercise B.5-3), # of internal nodes = # of leaves-1.

So, $n = \# \text{ of nodes} = \# \text{ of leaves} + \# \text{ of internal nodes} = 2 \# \text{ of leaves} - 1$.

Thus, # of leaves = $\frac{(n+1)}{2} \leq \left\lceil \frac{n}{2} \right\rceil$. (The latter equality holds because n is odd.)

- If n is even, then x is odd, and some leaf doesn't have a sibling. If we gave it a sibling, we would have $n + 1$ nodes, where $n + 1$ is odd, so the case we analyzed above would apply. Observe that we would also increase the number of leaves by 1, since we added a node to a parent that already had a child. By the odd-node case above, #

$$\text{of leaves} + 1 = \frac{(n + 1)}{2} \leq \left\lceil \frac{n}{2} \right\rceil + 1. \text{ (The latter equality holds because } n \text{ is even.)}$$

- In either case, # of leaves $\leq \left\lceil \frac{n}{2} \right\rceil$.

Inductive step: Show that if it's true for height $h - 1$, it's true for h .

Let n_h be the number of nodes at height h in the n -node tree T .

Consider the tree T' formed by removing the leaves of T . It has $n' = n - n_0$

nodes. We know from the base case that $n_0 = \left\lceil \frac{n}{2} \right\rceil$, so

$$n' = n - n_0 \leq n - \left\lceil \frac{n}{2} \right\rceil = \left\lfloor \frac{n}{2} \right\rfloor.$$

- Note that the nodes at height h in T would be at height $h - 1$ in T' .
Letting n'_{h-1} denote the number of nodes at height $h - 1$ in T' , we have
 $n_h = n'_{h-1}$

By induction, we can bound n'_{h-1} :

$$n_h = n'_{h-1} \leq \left\lceil \frac{n'}{2^h} \right\rceil \leq \left\lceil \frac{\left\lfloor \frac{n}{2} \right\rfloor}{2^h} \right\rceil \leq \left\lceil \frac{\frac{n}{2}}{2^h} \right\rceil = \left\lceil \frac{n}{2^{h+1}} \right\rceil$$



Analysis

- Analysis of BUILD-MAX-HEAP

- **Simple bound:** $O(n)$ calls to MAX-HEAPIFY, each of which takes $O(\lg n)$ time $\rightarrow O(n \lg n)$.

- **Tighter analysis:** Have $\leq \lceil n/2^{h+1} \rceil$ nodes of height h , and height of heap is $\lceil \lg n \rceil$.

- The Time required by MAX-HEAPIFY when called on a node of height h is $O(h)$, so the total cost of BUILD-MAX-HEAP is

$$\sum_{h=0}^{\lceil \lg n \rceil} \left\lceil \frac{n}{2^{h+1}} \right\rceil O(h) = O\left(n \sum_{h=0}^{\lceil \lg n \rceil} \frac{h}{2^h} \right)$$

- Evaluate the last summation by substituting $x = 1/2$ in the formula $\sum_{k=0}^{\infty} kx^k = \frac{x}{(1-x)^2}$ for $|x| < 1$.

- Thus, the running time of BUILD-MAX-HEAP is $O(n)$.

$$\sum_{h=0}^{\infty} \frac{h}{2^h} = \frac{1/2}{(1 - 1/2)^2} = 2$$

Building a min-heap

- Building a min-heap from an unordered array can be done by calling **MIN-HEAPIFY** instead of MAX-HEAPIFY, also taking linear time.








The heapsort algorithm

HEAPSORT(A, n)

```
1 BUILD-MAX-HEAP( $A, n$ )
2 for  $i = n$  downto 2 do
3     exchange  $A[1] \leftrightarrow A[i]$ 
4     MAX-HEAPIFY( $A, 1, i - 1$ )
```



Example: Heapsort Algorithm



A

1	2	3	4	7
---	---	---	---	---

Analysis

- Analysis of heapsort
 - BUILD-MAX-HEAP: $O(n)$
 - for loop: $n - 1$ times
 - Exchange elements: $O(1)$
 - MAX-HEAPIFY: $O(\lg n)$
- Total time: $O(n \lg n)$

- The worst-case running time of HEAPSORT is $\Omega(n \lg n)$
 - Whenever we have an array that is already sorted, we take linear time to convert it to a max-heap and then $n \lg n$ time to sort it.
- When all elements are distinct, the best-case running time of HEAPSORT is $\Omega(n \lg n)$
 - T. I. Fenner and A. M. Frieze, “On the Best Case of Heapsort”

Analysis

- Though heapsort is a great algorithm, it is usually not quite as fast as quicksort for large n .
- On the other hand, unlike quicksort, its performance is guaranteed.



Priority queue

- Heap implementation of priority queue
 - Max-priority queues are implemented with max-heaps. Min-priority queues are implemented with min-heaps similarly.
- Max Priority Queues
 - Maintains a dynamic set of S of elements.
 - Each with an associated value called a key.
 - Max-priority queue supports dynamic-set operations:
 - **INSERT(S, x)**: inserts element x into set S .
 - **MAXIMUM(S)**: returns elements of S with largest key.
 - **EXTRACT-MAX(S)**: removes and returns element of S with largest key.
 - **INCREASE-KEY(S, x, k)**: increases value of element x 's key to k . Assume $k \geq x$'s current key value.

- Min-priority queue supports similar operation:
 - $\text{INSERT}(S, x)$: inserts element x into set S .
 - $\text{MINIMUM}(S)$: returns element of S with smallest key.
 - $\text{EXTRACT-MIN}(S)$: removes and returns element of S with smallest key.
 - $\text{DECREASE-KEY}(S, x, k)$: decreases value of element x 's key to k . Assume $k \leq x$'s current key value.


- Finding the maximum element

- Getting the maximum element is easy: it's the root.

HEAP-MAXIMUM(*A*)

1 return *A*[1]

- Time: $\Theta(1)$




HEAP-MAXIMUM(*A*) returns 7

- Extracting max element

- Given the array A :
 - Make sure heap is not empty
 - Make a copy of the maximum element (the root).
 - Make the last node in the tree the new root.
 - Re-heapify the heap, with one fewer node.
 - Return the copy of the maximum element.

HEAP-EXTRACT-MAX(A, n)

```
1 if  $n < 1$  then  
2   error "heap underflow"  
3  $max \leftarrow A[1]$   
4  $A[1] \leftarrow A[n]$   
5 MAX-HEAPIFY( $A, 1, n - 1$ ) ▶remakes heap  
6 return  $max$ 
```



- **Analysis:** constant time assignments plus time for MAX-HEAPIFY.
- **Time:** $O(\lg n)$


Increasing key value

- Given set S , element x and new key value k :
 - Make sure $k \geq x$'s current key.
 - Update x 's key value to k .
 - Traverse the tree upward comparing x to its parent and swapping keys if necessary, until x 's key is smaller than parent's key.



HEAP-INCREASE-KEY(A, i, key)

```
1 if  $key < A[i]$  then  
2   error "new key is smaller than current key"  
3  $A[i] \leftarrow key$   
4 while  $i > 1$  and  $A[PARENT(i)] < A[i]$  do  
5   exchange  $A[i] \leftrightarrow A[PARENT(i)]$   
6    $i \leftarrow PARENT(i)$ 
```








HEAP-INCREASE-KEY(A , 5, 15)





$A[i] \leftarrow \text{key}$



$A[\text{PARANT}(i)] < A[i]$




$A[\text{PARANT}(i)] \leftrightarrow A[i]$





$i \leftarrow \text{PARANT}(i)$







$$A[\text{PARANT}(i)] < A[i]$$



$$A[\text{PARANT}(i)] \leftrightarrow A[i]$$



$$i \leftarrow \text{PARANT}(i)$$




- **Analysis:** Upward path from node i has length $O(\lg n)$ in an n -element heap.
- **Time:** $O(\lg n)$

Inserting into the heap


- Given a key k to insert into the heap:
 - Insert a new node in the key last position in the tree with key $-\infty$
 - Increase the $-\infty$ key to k using the HEAP-INCREASE-KEY procedure defined above.

MAX-HEAP-INSERT(A, key, n)

- 1 $A[i + 1] \leftarrow -\infty$
- 2 HEAP-INCREASE-KEY($A, n + 1, key$)



- **Analysis:** constant time assignments plus time for HEAP-INCREASE-KEY
- **Time:** $O(\lg n)$
- **Min-priority queue** operations are implemented similarly with min-heaps.



藏行顯光 成就共好

Achieve Securely
Prosper Mutually



國立成功大學 九十週年
90th Anniversary of NCKU



國立成功大學
National Cheng Kung University